

```
In [ ]: dir='/kaggle/input/gender-classification-dataset/gender_classification_v7.csv'

In [ ]: import pandas as pd
import numpy as np
import missingno as msng
import seaborn as sns
import matplotlib.pyplot as plt
import scipy.stats as stats

In [ ]: df = pd.read_csv(dir)
df.head()
```

Out[ ]:	long_hair	forehead_width_cm	forehead_height_cm	nose_wide	nose_long	lips_thin	distance_nose_to_lip_long	gender
0	1	11.8	6.1	1	0	1	1	Male
1	0	14.0	5.4	0	0	1	0	Female
2	0	11.8	6.3	1	1	1	1	Male
3	0	14.4	6.1	0	1	1	1	Male
4	1	13.5	5.9	0	0	0	0	Female

**Problem Introduction:**

**Objective:** Predict gender based on facial attributes.

**Dataset Features:**

- 'long\_hair': Presence of long hair (1: Yes, 0: No).
- 'forehead\_width\_cm' and 'forehead\_height\_cm': Dimensions of the forehead.
- 'nose\_wide' and 'nose\_long': Width and length of the nose.
- 'lips\_thin': Thinness of the lips.
- 'distance\_nose\_to\_lip\_long': Length of the distance from nose to lip.

**Target Variable:** 'gender' (Male or Female).

**Data Set Description:**

```
In [ ]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5001 entries, 0 to 5000
Data columns (total 8 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   long_hair                             5001 non-null   int64
1   forehead_width_cm                     5001 non-null   float64
2   forehead_height_cm                    5001 non-null   float64
3   nose_wide                             5001 non-null   int64
4   nose_long                             5001 non-null   int64
5   lips_thin                             5001 non-null   int64
6   distance_nose_to_lip_long              5001 non-null   int64
7   gender                                5001 non-null   object
dtypes: float64(2), int64(5), object(1)
memory usage: 312.7+ KB

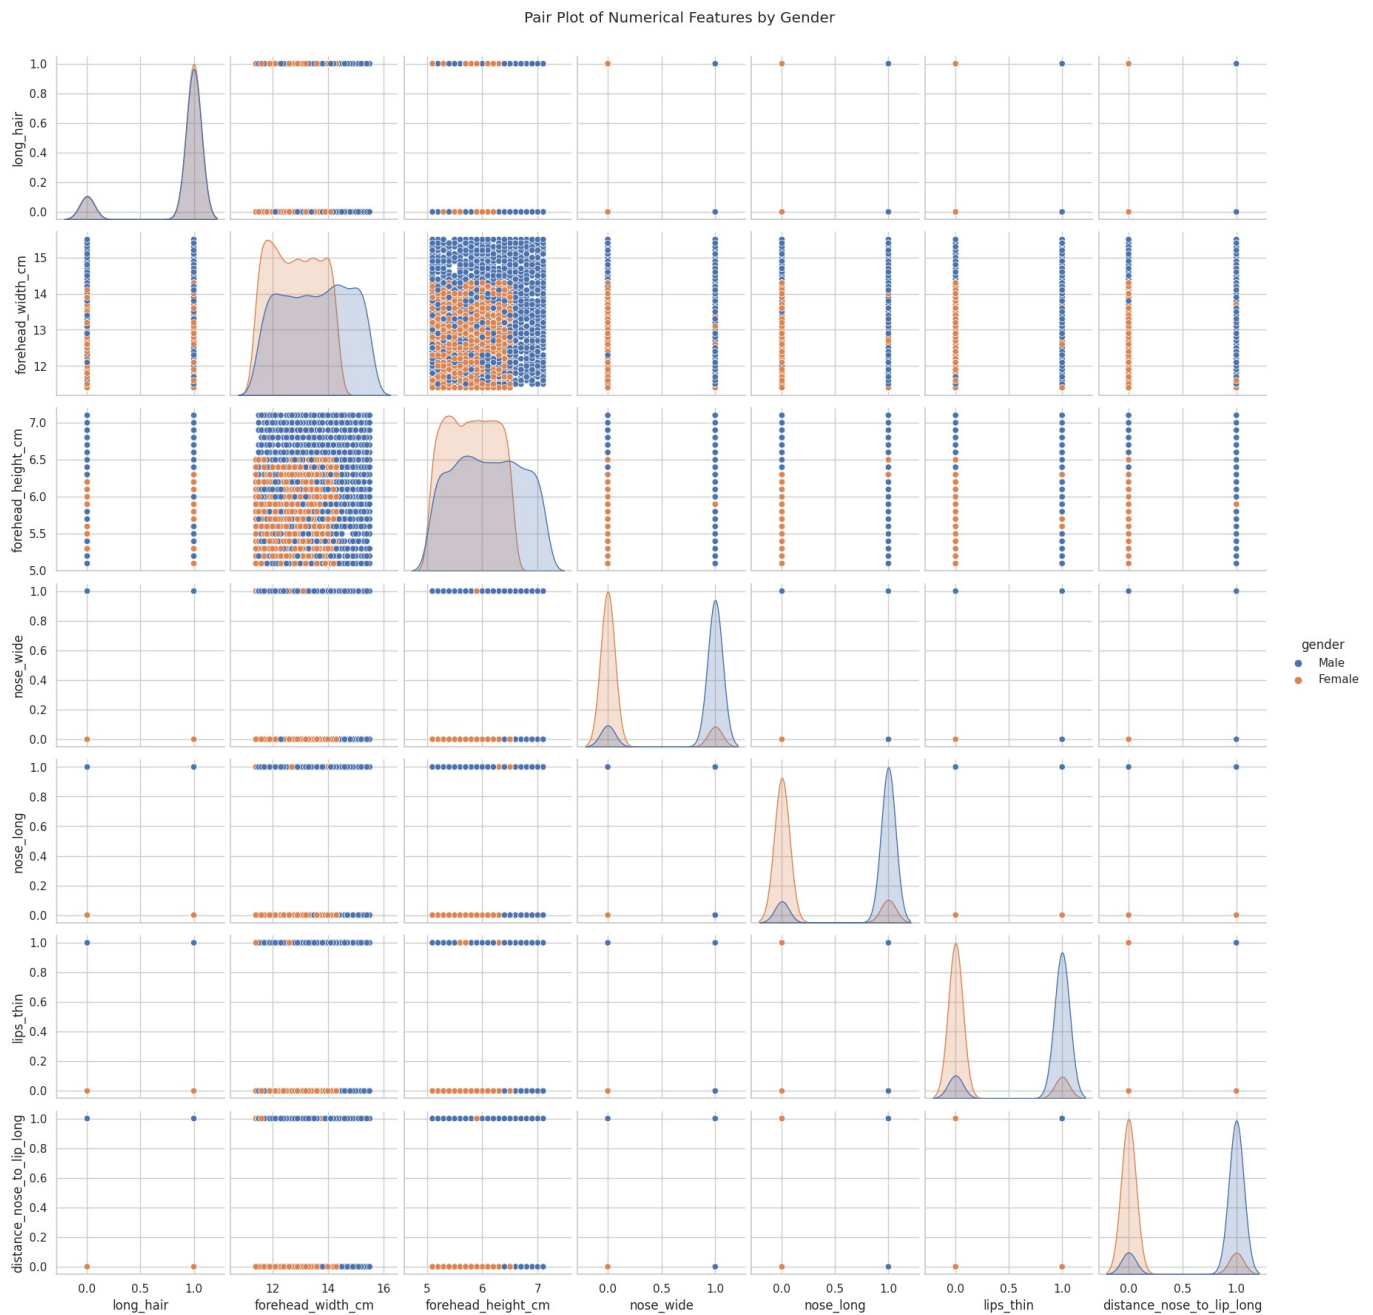
In [ ]: df.describe()
```

Out[ ]:	long_hair	forehead_width_cm	forehead_height_cm	nose_wide	nose_long	lips_thin	distance_nose_to_lip_long
count	5001.000000	5001.000000	5001.000000	5001.000000	5001.000000	5001.000000	5001.000000
mean	0.869626	13.181484	5.946311	0.493901	0.507898	0.493101	0.498900
std	0.336748	1.107128	0.541268	0.500013	0.499988	0.500002	0.500049
min	0.000000	11.400000	5.100000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	12.200000	5.500000	0.000000	0.000000	0.000000	0.000000
50%	1.000000	13.100000	5.900000	0.000000	1.000000	0.000000	0.000000
75%	1.000000	14.000000	6.400000	1.000000	1.000000	1.000000	1.000000
max	1.000000	15.500000	7.100000	1.000000	1.000000	1.000000	1.000000

## Visualizations:


```
In [ ]: # Set the style for seaborn
sns.set(style="whitegrid")
```

```
In [ ]: #Pair Plot for all numerical features
sns.pairplot(df, hue='gender', diag_kind='kde')
plt.suptitle('Pair Plot of Numerical Features by Gender', y=1.02)
plt.show()
```




```
In [ ]: # Boxplot for numerical features by gender
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(12, 6))
sns.boxplot(x="gender", y="forehead_width_cm", data=df, ax=axes[0], palette={"Male": "blue", "Female": "pink"})
sns.boxplot(x="gender", y="forehead_height_cm", data=df, ax=axes[1], palette={"Male": "blue", "Female": "pink"})
plt.suptitle("Boxplots of Numerical Features by Gender", y=1.02)
plt.tight_layout()
plt.show()
```

Boxplots of Numerical Features by Gender




```
In [ ]: # Countplot for categorical features
fig, axes = plt.subplots(nrows=2, ncols=2, figsize=(12, 10))
sns.countplot(x="long_hair", hue="gender", data=df, ax=axes[0, 0], palette={"Male": "blue", "Female": "pink"})
sns.countplot(x="nose_wide", hue="gender", data=df, ax=axes[0, 1], palette={"Male": "blue", "Female": "pink"})
sns.countplot(x="nose_long", hue="gender", data=df, ax=axes[1, 0], palette={"Male": "blue", "Female": "pink"})
sns.countplot(x="lips_thin", hue="gender", data=df, ax=axes[1, 1], palette={"Male": "blue", "Female": "pink"})
plt.suptitle("Countplots of Categorical Features by Gender", y=1.02)
plt.tight_layout()
plt.show()
```


Countplots of Categorical Features by Gender



```
In [ ]: #Violin Plot for 'forehead_height_cm' by gender
sns.violinplot(x='gender', y='forehead_height_cm', data=df, palette='pastel')
plt.title('Distribution of Forehead Height by Gender')
plt.show()
```




```
In [ ]: # Scatter Plot for 'forehead_width_cm' vs 'forehead_height_cm'
sns.scatterplot(x='forehead_width_cm', y='forehead_height_cm', hue='gender', data=df, palette='Set1')
plt.title('Scatter Plot of Forehead Dimensions by Gender')
plt.show()
```

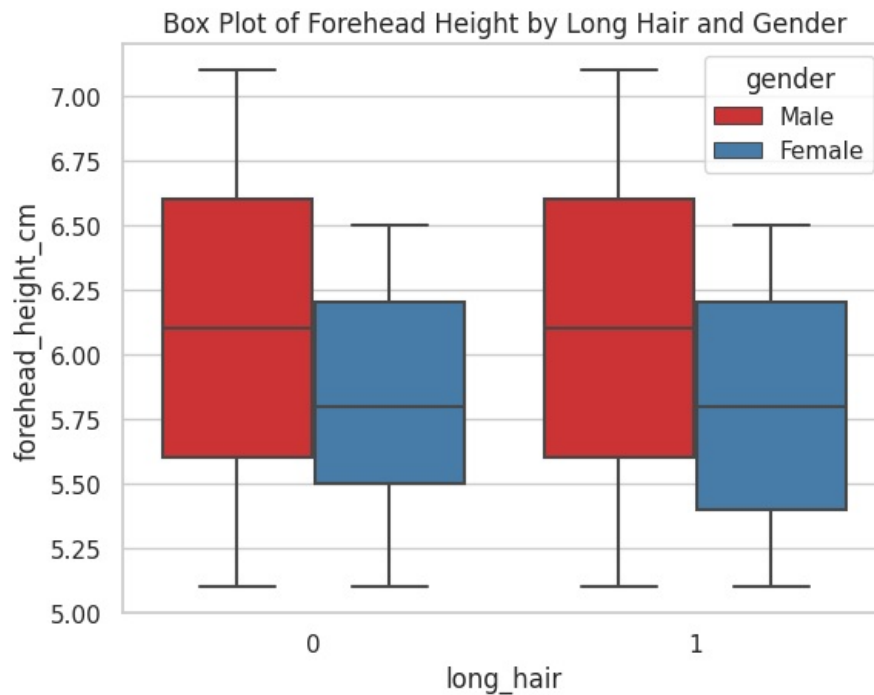


```
In [ ]: #Swarm Plot for 'long_hair' and 'forehead_width_cm' by gender
sns.swarmplot(x='long_hair', y='forehead_width_cm', hue='gender', data=df, palette='Set2')
plt.title('Swarm Plot of Long Hair and Forehead Width by Gender')
plt.show()
```

/opt/conda/lib/python3.10/site-packages/seaborn/categorical.py:3544: UserWarning: 52.8% of the points cannot be placed; you may want to decrease the size of the markers or use stripplot.  
 warnings.warn(msg, UserWarning)  
/opt/conda/lib/python3.10/site-packages/seaborn/categorical.py:3544: UserWarning: 73.9% of the points cannot be placed; you may want to decrease the size of the markers or use stripplot.  
 warnings.warn(msg, UserWarning)



```
In [ ]: #Box Plot for 'forehead_height_cm' by 'long_hair' and 'gender'
sns.boxplot(x='long_hair', y='forehead_height_cm', hue='gender', data=df, palette='Set1')
plt.title('Box Plot of Forehead Height by Long Hair and Gender')
plt.show()
```



### Missing Values Treatment

```
In [ ]: df.isnull().sum()
```

```
Out[ ]: long_hair          0
forehead_width_cm       0
forehead_height_cm      0
nose_wide               0
nose_long              0
lips_thin              0
distance_nose_to_lip_long 0
gender                 0
dtype: int64
```

\*\*\*\*\*Binning\*\*\*\*\*

```
In [ ]: # Example
# df['forehead_width_bin'] = pd.cut(df['forehead_width_cm'], bins=5, labels=False)
```

### Data Analysis

```
In [ ]: numerical_columns = df.select_dtypes(include=['float64', 'int64']).columns
```

```
In [ ]: summary_statistics = df[numerical_columns].agg(['min', 'max', 'mean', 'var', 'std', 'skew', 'kurt'])
```

```
In [ ]: print("Summary Statistics:")
print(summary_statistics)
```

```
Summary Statistics:
      long_hair  forehead_width_cm  forehead_height_cm  nose_wide  nose_long  \
min    0.000000         11.400000         5.100000    0.000000    0.000000
max    1.000000         15.500000         7.100000    1.000000    1.000000
mean   0.869626         13.181484         5.946311    0.493901    0.507898
var    0.113399         1.225733         0.292971    0.250013    0.249988
std    0.336748         1.107128         0.541268    0.500013    0.499988
skew   -2.196146         0.242242         0.250739    0.024404   -0.031607
kurt    2.824187        -0.930596        -0.848889   -2.000205   -1.999801

      lips_thin  distance_nose_to_lip_long
min    0.000000         0.000000
max    1.000000         1.000000
mean   0.493101         0.498900
var    0.250002         0.250049
std    0.500002         0.500049
skew   0.027605         0.004400
kurt   -2.000038        -2.000781
```

### Data Analysis

```
In [ ]: from scipy.stats import chi2_contingency, ttest_ind, f_oneway
```

```
In [ ]: # Covariance Matrix
covariance_matrix = df[numerical_columns].cov()
covariance_matrix
```

```
Out[ ]:
```

	long_hair	forehead_width_cm	forehead_height_cm	nose_wide	nose_long	lips_thin	distance_nose_to_lip_long
long_hair	0.113399	-0.002435	-0.003141	0.000205	0.002430	0.001900	-0.004343
forehead_width_cm	-0.002435	1.225733	0.053092	0.139307	0.142466	0.143132	0.139140
forehead_height_cm	-0.003141	0.053092	0.292971	0.057282	0.052534	0.055600	0.058271
nose_wide	0.000205	0.139307	0.057282	0.250013	0.141298	0.139408	0.142343
nose_long	0.002430	0.142466	0.052534	0.141298	0.249988	0.140304	0.139959
lips_thin	0.001900	0.143132	0.055600	0.139408	0.140304	0.250002	0.141342
distance_nose_to_lip_long	-0.004343	0.139140	0.058271	0.142343	0.139959	0.141342	0.250002

```
In [ ]: # Correlation
correlation_matrix = df[numerical_columns].corr()
correlation_matrix
```

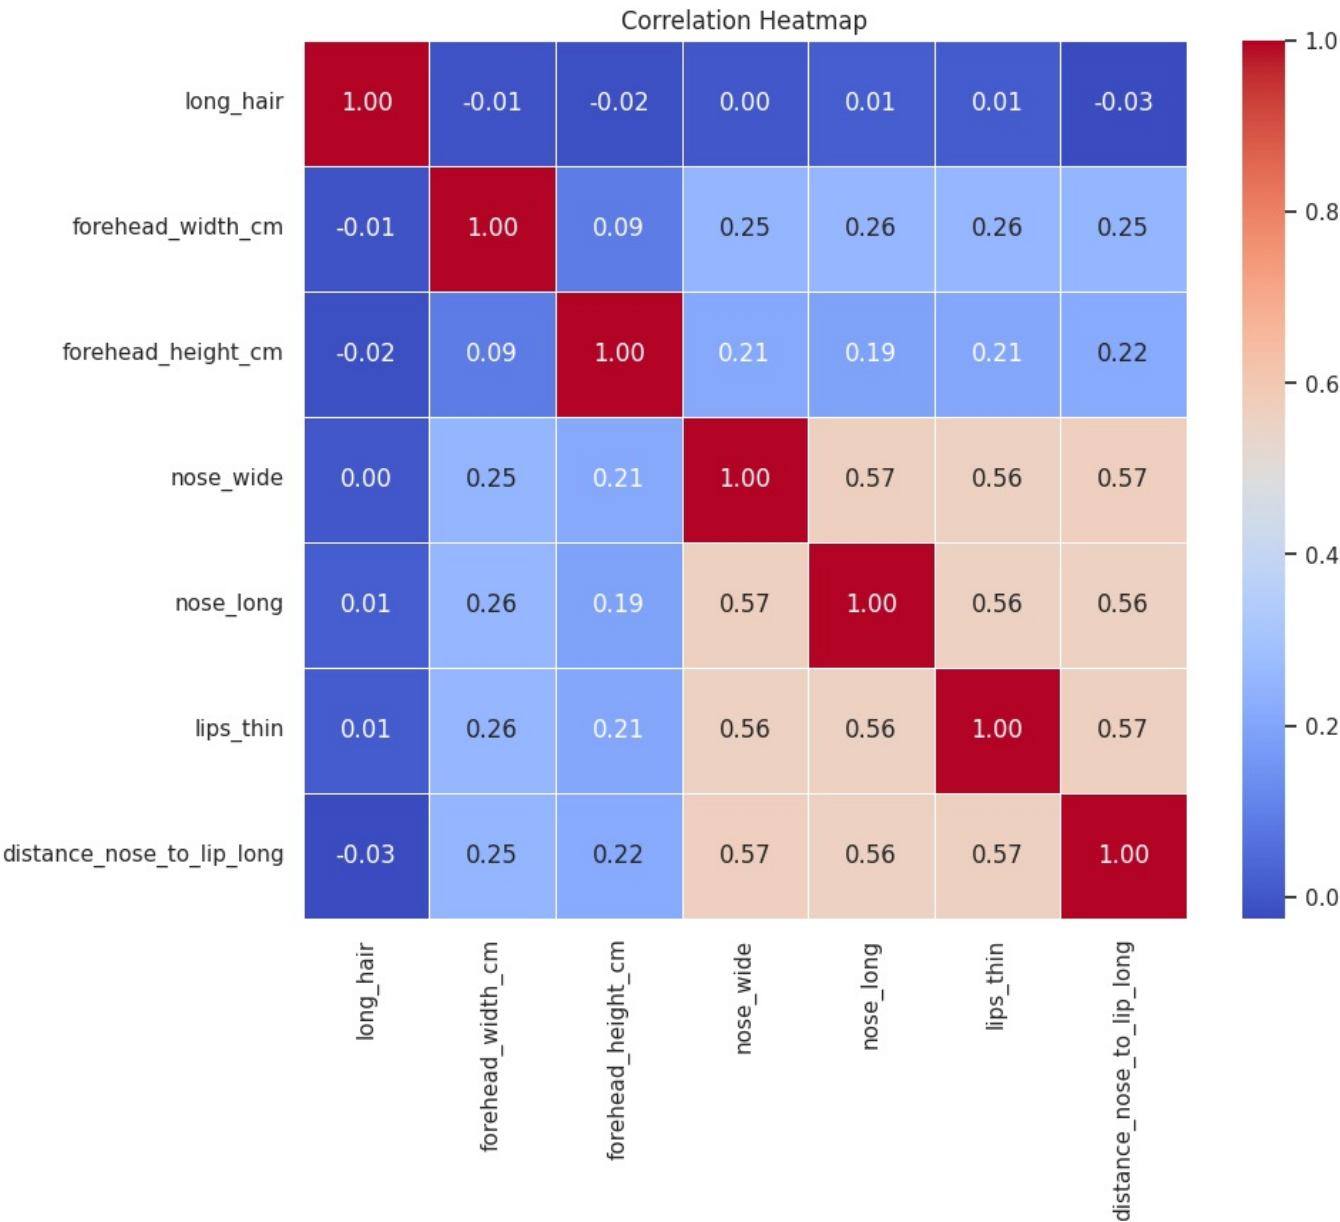


Out [ ]:

	long_hair	forehead_width_cm	forehead_height_cm	nose_wide	nose_long	lips_thin	distance_nose_to_lip_long
long_hair	1.000000	-0.006530	-0.017233	0.001216	0.014432	0.011287	-0
forehead_width_cm	-0.006530	1.000000	0.088596	0.251648	0.257368	0.258564	0
forehead_height_cm	-0.017233	0.088596	1.000000	0.211655	0.194120	0.205441	0
nose_wide	0.001216	0.251648	0.211655	1.000000	0.565192	0.557615	0
nose_long	0.014432	0.257368	0.194120	0.565192	1.000000	0.561229	0
lips_thin	0.011287	0.258564	0.205441	0.557615	0.561229	1.000000	0
distance_nose_to_lip_long	-0.025794	0.251328	0.215292	0.569303	0.559794	0.565312	1

In [ ]:

```
# Heatmap for Correlation Matrix
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f", linewidths=.5)
plt.title('Correlation Heatmap')
plt.show()
```



In [ ]:

```
# Chi-square Test
contingency_table = pd.crosstab(df['long_hair'], df['gender'])
chi2_stat, p_value, _, _ = chi2_contingency(contingency_table)
print(f"Chi-square Statistic: {chi2_stat}\nP-value: {p_value}")

Chi-square Statistic: 0.517551479459157
P-value: 0.47188802409575925

# Z-test or t-test
male_forehead = df[df['gender'] == 'Male']['forehead_width_cm']
female_forehead = df[df['gender'] == 'Female']['forehead_width_cm']
t_stat, p_value_ttest = ttest_ind(male_forehead, female_forehead)
print(f"T-test Statistic: {t_stat}\nP-value: {p_value_ttest}")
```

T-test Statistic: 25.06432518851344  
P-value: 1.063197016698364e-130

```
In [ ]: # ANOVA (Example for 'forehead_width_cm' across different 'long_hair' categories)
anova_result = f_oneway(df['forehead_width_cm'][df['long_hair'] == 0],
                        df['forehead_width_cm'][df['long_hair'] == 1])
print(f"ANOVA F-statistic: {anova_result.statistic}\nP-value: {anova_result.pvalue}")
```

ANOVA F-statistic: 0.21316903345667512  
P-value: 0.6443148961556857

---

## Feature Reduction

### Linear Discriminant Analysis (LDA):

```
In [ ]: from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report
```

```
In [ ]: X = df.drop('gender', axis=1)
y = df['gender']
```

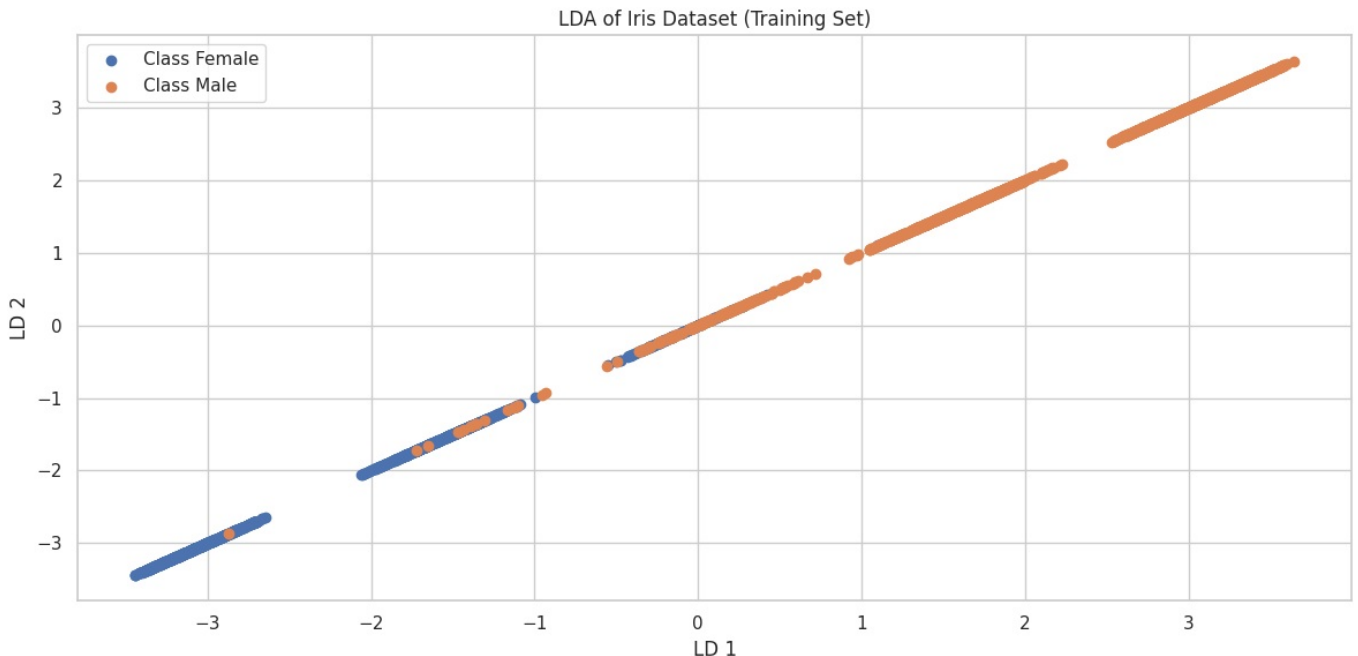
```
In [ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [ ]: n_components = min(X.shape[1], len(set(y)) - 1)
lda = LinearDiscriminantAnalysis(n_components=n_components)
```

```
In [ ]: X_lda_train = lda.fit_transform(X_train, y_train)
```

```
In [ ]: plt.figure(figsize=(12, 6))
for label in np.unique(y_train):
    plt.scatter(X_lda_train[y_train == label, 0], X_lda_train[y_train == label, 1], label=f'Class {label}')

plt.title('LDA of Iris Dataset (Training Set)')
plt.xlabel('LD 1')
plt.ylabel('LD 2')
plt.legend()
plt.tight_layout()
plt.show()
```



### Principal Component Analysis (PCA):

```
In [ ]: from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler, LabelEncoder
```

```
In [ ]: X = df.drop('gender', axis=1)
y = df['gender']
```

```
In [ ]: label_encoder = LabelEncoder()
y_encoded = label_encoder.fit_transform(y)
```

```
In [ ]: scaler = StandardScaler()
```



```
X_scaled = scaler.fit_transform(X)
```

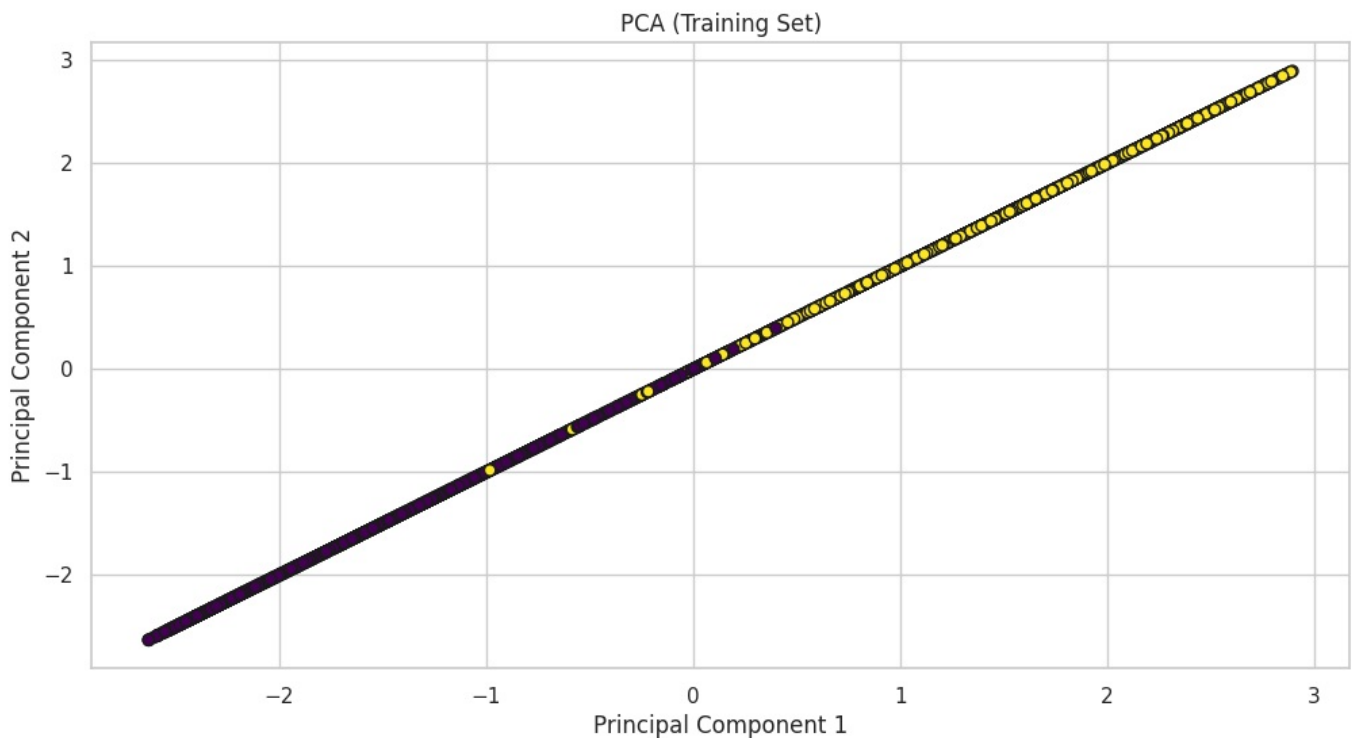
```
In [ ]: n_components = min(X.shape[1], len(set(y)) - 1)
pca = PCA(n_components=n_components)
```

```
In [ ]: X_train_pca = pca.fit_transform(X_scaled)
```

```
In [ ]: print("Original shape:", X.shape)
print("Transformed shape:", X_train_pca.shape)
```

Original shape: (5001, 7)  
Transformed shape: (5001, 1)

```
In [ ]: plt.figure(figsize=(12, 6))
# Plotting the first principal component against the second principal component
plt.scatter(X_train_pca[:, 0], X_train_pca[:,1], c=y_encoded, cmap='viridis', edgecolor='k')
plt.title('PCA (Training Set)')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.show()
```



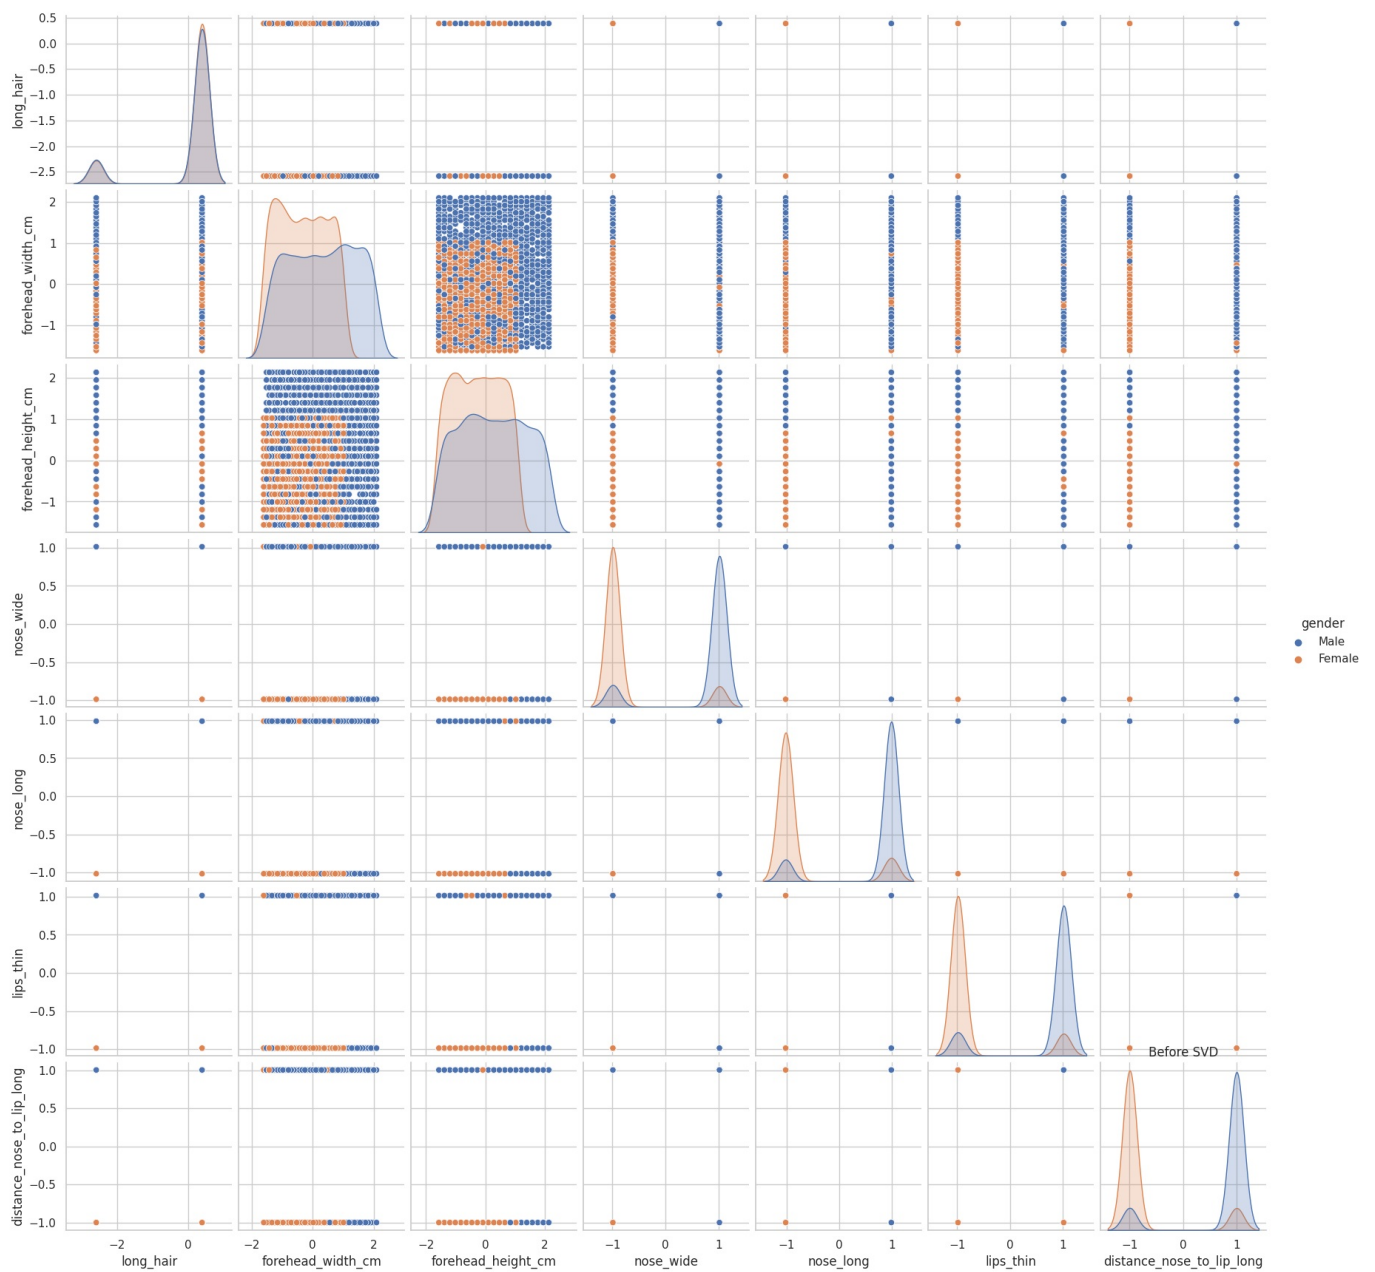
### Singular Value Decomposition (SVD)

```
In [ ]: from sklearn.decomposition import TruncatedSVD
```

```
In [ ]: X = df.drop('gender', axis=1)
y = df['gender']
```

```
In [ ]: scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

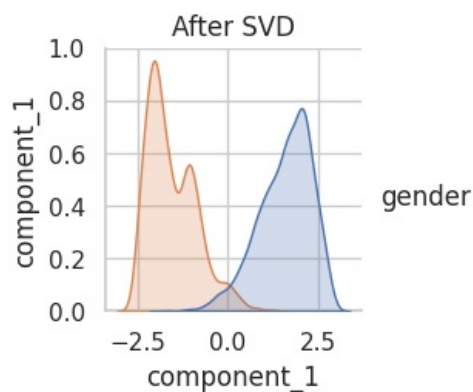
```
In [ ]: # Before SVD Visualization
sns.pairplot(pd.concat([pd.DataFrame(X_scaled, columns=X.columns), y], axis=1), hue='gender')
plt.title('Before SVD')
plt.show()
```



```
In [ ]: n_components = min(X.shape[1], len(set(y)) - 1)
svd = TruncatedSVD(n_components=n_components)
X_svd = svd.fit_transform(X_scaled)
```

```
In [ ]: column_names = [f'component_{i+1}' for i in range(n_components)]
df_svd = pd.DataFrame(data=X_svd, columns=column_names)
df_svd['gender'] = y
```

```
In [ ]: sns.pairplot(df_svd, hue='gender')
plt.title('After SVD')
plt.show()
```



## Model Implementations

Naive Bayes (Gaussian Naive Bayes):

```
In [ ]: classification_reports = {}
```

```
In [ ]: from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report
from sklearn.model_selection import cross_val_score, cross_val_predict
from sklearn.metrics import confusion_matrix, roc_curve, auc
from sklearn.preprocessing import LabelBinarizer
from sklearn.multiclass import OneVsRestClassifier
```

```
In [ ]: X = df.drop('gender', axis=1)
y = df['gender']
```

```
In [ ]: # Convert categorical labels to binary labels
lb = LabelBinarizer()
y_bin = lb.fit_transform(y)
```

```
In [ ]: # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y_bin, test_size=0.2, random_state=42)
```

```
In [ ]: # Initialize Gaussian Naive Bayes model
nb_model = OneVsRestClassifier(GaussianNB())
```

```
In [ ]: # Fit the model
nb_model.fit(X_train, y_train)
```

```
Out[ ]: ▸ OneVsRestClassifier
▸ estimator: GaussianNB
    ▸ GaussianNB
```

```
In [ ]: # Predict on the test set
y_pred_nb = nb_model.predict(X_test)
```

```
In [ ]: # Evaluate the model
accuracy_nb = accuracy_score(y_test, y_pred_nb)
print("Naive Bayes Accuracy:", accuracy_nb)
```

Naive Bayes Accuracy: 0.964035964035964

```
In [ ]: # Additional evaluation metrics
print("Classification Report:")
print(classification_report(y_test, y_pred_nb))
```

Classification Report:

	precision	recall	f1-score	support
0	0.96	0.97	0.96	502
1	0.97	0.96	0.96	499
accuracy			0.96	1001
macro avg	0.96	0.96	0.96	1001
weighted avg	0.96	0.96	0.96	1001

```
In [ ]: classification_reports['Naive Bayes'] = classification_report(y_test, y_pred_nb)
```

```
In [ ]: # K-fold cross-validation and average accuracy
cv_accuracy = cross_val_score(nb_model, X, y_bin, cv=10, scoring='accuracy')
avg_cv_accuracy = cv_accuracy.mean()
print("Average Cross-Validation Accuracy:", avg_cv_accuracy)
```

Average Cross-Validation Accuracy: 0.9702091816367264

```
In [ ]: # Confusion Matrix
conf_matrix = confusion_matrix(y_test, y_pred_nb)
print("Confusion Matrix:")
print(conf_matrix)
```

Confusion Matrix:  
[[487 15]  
 [ 21 478]]

```
In [ ]: # Accuracy, Error rate, Precision, Recall, F-measure
tp, fp, fn, tn = conf_matrix.ravel()
accuracy = (tp + tn) / (tp + fp + fn + tn)
error_rate = 1 - accuracy
precision = tp / (tp + fp)
recall = tp / (tp + fn)
```

```
f_measure = 2 * (precision * recall) / (precision + recall)
```

```
print("Accuracy:", accuracy)
print("Error Rate:", error_rate)
print("Precision:", precision)
print("Recall:", recall)
print("F-measure:", f_measure)
```

Accuracy: 0.964035964035964  
Error Rate: 0.03596403596403597  
Precision: 0.9701195219123506  
Recall: 0.9586614173228346  
F-measure: 0.9643564356435643

In [ ]: # ROC Curve


```
y_scores = cross_val_predict(nb_model, X, y_bin.ravel(), cv=10, method="predict_proba")[:, 1]
fpr, tpr, thresholds = roc_curve(y_bin.ravel(), y_scores)
roc_auc = auc(fpr, tpr)

print("ROC AUC:", roc_auc)
```

ROC AUC: 0.9966251099560175

In [ ]: # Visualize ROC Curve

```
plt.figure(figsize=(8, 8))
plt.plot(fpr, tpr, label=f'ROC Curve (AUC = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], 'k--', label='Random Guessing')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend()
plt.show()
```



In [ ]: # Interpret results of the confusion matrix

```
if tp + fp == 0 or tp + fn == 0:
    print("Unable to determine if the model is overfitting or underfitting.")
else:
    precision_positive = tp / (tp + fp)
    recall_positive = tp / (tp + fn)
    specificity = tn / (tn + fp)
    sensitivity = recall_positive
    balance = 1 - abs(1 - (precision_positive + sensitivity) / 2)
```

```
print("Precision (Positive):", precision_positive)
print("Recall (Positive):", recall_positive)
print("Specificity:", specificity)
print("Sensitivity:", sensitivity)
print("Balance:", balance)

if balance > 0.99:
    print("The model may be overfitting.")
elif balance < 0.05:
    print("The model may be underfitting.")
else:
    print("The model is reasonably balanced.")
```

```
Precision (Positive): 0.9701195219123506
Recall (Positive): 0.9586614173228346
Specificity: 0.9695740365111561
Sensitivity: 0.9586614173228346
Balance: 0.9643904696175927
The model is reasonably balanced.
```

#### **Bayesian Belief Network:**

```
In [ ]: !pip install pgmpy
```

Collecting pgmpy

Obtaining dependency information for pgmpy from <https://files.pythonhosted.org/packages/eb/9a/2fcb6fd998a016cef29ca3eab30b98b6c232b6e9a0444df07f0ad47f8d/pgmpy-0.1.24-py3-none-any.whl.metadata>

Downloading pgmpy-0.1.24-py3-none-any.whl.metadata (6.3 kB)

Requirement already satisfied: networkx in /opt/conda/lib/python3.10/site-packages (from pgmpy) (3.1)

Requirement already satisfied: numpy in /opt/conda/lib/python3.10/site-packages (from pgmpy) (1.24.3)

Requirement already satisfied: scipy in /opt/conda/lib/python3.10/site-packages (from pgmpy) (1.11.4)

Requirement already satisfied: scikit-learn in /opt/conda/lib/python3.10/site-packages (from pgmpy) (1.2.2)

Requirement already satisfied: pandas in /opt/conda/lib/python3.10/site-packages (from pgmpy) (2.0.3)

Requirement already satisfied: pyparsing in /opt/conda/lib/python3.10/site-packages (from pgmpy) (3.0.9)

Requirement already satisfied: torch in /opt/conda/lib/python3.10/site-packages (from pgmpy) (2.0.0+cpu)

Requirement already satisfied: statsmodels in /opt/conda/lib/python3.10/site-packages (from pgmpy) (0.14.0)

Requirement already satisfied: tqdm in /opt/conda/lib/python3.10/site-packages (from pgmpy) (4.66.1)

Requirement already satisfied: joblib in /opt/conda/lib/python3.10/site-packages (from pgmpy) (1.3.2)

Requirement already satisfied: opt-einsum in /opt/conda/lib/python3.10/site-packages (from pgmpy) (3.3.0)

Requirement already satisfied: python-dateutil<=2.8.2 in /opt/conda/lib/python3.10/site-packages (from pandas->pgmpy) (2.8.2)

Requirement already satisfied: pytz<=2020.1 in /opt/conda/lib/python3.10/site-packages (from pandas->pgmpy) (2023.3)

Requirement already satisfied: tzdata<=2022.1 in /opt/conda/lib/python3.10/site-packages (from pandas->pgmpy) (2023.3)

Requirement already satisfied: threadpoolctl<=2.0.0 in /opt/conda/lib/python3.10/site-packages (from scikit-learn->pgmpy) (3.2.0)

Requirement already satisfied: patsy<=0.5.2 in /opt/conda/lib/python3.10/site-packages (from statsmodels->pgmpy) (0.5.3)

Requirement already satisfied: packaging<=21.3 in /opt/conda/lib/python3.10/site-packages (from statsmodels->pgmpy) (21.3)

Requirement already satisfied: filelock in /opt/conda/lib/python3.10/site-packages (from torch->pgmpy) (3.12.2)

Requirement already satisfied: typing-extensions in /opt/conda/lib/python3.10/site-packages (from torch->pgmpy) (4.5.0)

Requirement already satisfied: sympy in /opt/conda/lib/python3.10/site-packages (from torch->pgmpy) (1.12)

Requirement already satisfied: jinja2 in /opt/conda/lib/python3.10/site-packages (from torch->pgmpy) (3.1.2)

Requirement already satisfied: six in /opt/conda/lib/python3.10/site-packages (from patsy<=0.5.2->statsmodels->pgmpy) (1.16.0)

Requirement already satisfied: MarkupSafe<=2.0 in /opt/conda/lib/python3.10/site-packages (from jinja2->torch->pgmpy) (2.1.3)

Requirement already satisfied: mpmath<=0.19 in /opt/conda/lib/python3.10/site-packages (from sympy->torch->pgmpy) (1.3.0)

Downloading pgmpy-0.1.24-py3-none-any.whl (2.0 MB)

2.0/2.0 MB 22.9 MB/s eta 0:00:00

Installing collected packages: pgmpy

Successfully installed pgmpy-0.1.24

```
In [ ]: from pgmpy.models import BayesianModel
        from pgmpy.estimators import MaximumLikelihoodEstimator
```

```
In [ ]: model_structure = [('long_hair', 'gender'), ('forehead_width_cm', 'gender'), ('forehead_height_cm', 'gender'),
                          ('nose_wide', 'gender'), ('nose_long', 'gender'), ('lips_thin', 'gender'),
                          ('distance_nose_to_lip_long', 'gender')]
```

```
In [ ]: # Create a BayesianModel object
        bayesian_model = BayesianModel(model_structure)
```



```
In [ ]: # Fit the model parameters using Maximum Likelihood Estimation
model = MaximumLikelihoodEstimator(bayesian_model, df)
```

```
In [ ]: # Get the CPDs (Conditional Probability Distributions)
cpds = model.get_parameters()
```

```
In [ ]: # Add CPDs to the Bayesian Model
bayesian_model.add_cpds(*cpds)
```

#### Decision Tree (Entropy and Error Estimation):

```
In [ ]: from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report
```

```
In [ ]: X = df.drop('gender', axis=1)
y = df['gender']
```

```
In [ ]: # Convert categorical labels to binary labels
lb = LabelBinarizer()
y_bin = lb.fit_transform(y)
```

```
In [ ]: X_train, X_test, y_train, y_test = train_test_split(X, y_bin, test_size=0.2, random_state=42)
```

```
In [ ]: # Initialize Decision Tree model with entropy criterion
dt_entropy_model = DecisionTreeClassifier(criterion='entropy')
```

```
In [ ]: dt_entropy_model.fit(X_train, y_train)
```

```
Out[ ]: ▾ DecisionTreeClassifier
DecisionTreeClassifier(criterion='entropy')
```

```
In [ ]: # Predict on the test set
y_pred_dt_entropy = dt_entropy_model.predict(X_test)
```

```
In [ ]: # Evaluate the model
accuracy_dt_entropy = accuracy_score(y_test, y_pred_dt_entropy)
print("Decision Tree (Entropy) Accuracy:", accuracy_dt_entropy)
```

Decision Tree (Entropy) Accuracy: 0.955044955044955

```
In [ ]: # Additional evaluation metrics
print("Classification Report:")
print(classification_report(y_test, y_pred_dt_entropy))
```

Classification Report:

	precision	recall	f1-score	support
0	0.96	0.95	0.96	502
1	0.95	0.96	0.96	499
accuracy			0.96	1001
macro avg	0.96	0.96	0.96	1001
weighted avg	0.96	0.96	0.96	1001

```
In [ ]: classification_reports['Desicion Tree Entropy'] = classification_report(y_test, y_pred_dt_entropy)
```

```
In [ ]: # K-fold cross-validation and average accuracy
cv_accuracy = cross_val_score(dt_entropy_model, X, y_bin, cv=10, scoring='accuracy')
avg_cv_accuracy = cv_accuracy.mean()
print("Average Cross-Validation Accuracy:", avg_cv_accuracy)
```

Average Cross-Validation Accuracy: 0.9638115768463074

```
In [ ]: # Confusion Matrix
conf_matrix = confusion_matrix(y_test, y_pred_dt_entropy)
print("Confusion Matrix:")
print(conf_matrix)
```

Confusion Matrix:  
[[478 24]  
 [ 21 478]]

```
In [ ]: # Accuracy, Error rate, Precision, Recall, F-measure
tp, fp, fn, tn = conf_matrix.ravel()
accuracy = (tp + tn) / (tp + fp + fn + tn)
error_rate = 1 - accuracy
precision = tp / (tp + fp)
```

```

recall = tp / (tp + fn)
f_measure = 2 * (precision * recall) / (precision + recall)

print("Accuracy:", accuracy)
print("Error Rate:", error_rate)
print("Precision:", precision)
print("Recall:", recall)
print("F-measure:", f_measure)

```

```

Accuracy: 0.955044955044955
Error Rate: 0.04495504495504499
Precision: 0.952191235059761
Recall: 0.9579158316633266
F-measure: 0.955044955044955

```

```

In [ ]: # ROC Curve
y_scores = cross_val_predict(dt_entropy_model, X, y_bin.ravel(), cv=10, method="predict_proba")[:, 1]
fpr, tpr, thresholds = roc_curve(y_bin.ravel(), y_scores)
roc_auc = auc(fpr, tpr)

print("ROC AUC:", roc_auc)


```

```
ROC AUC: 0.964481087564974
```

```

In [ ]: # Visualize ROC Curve
plt.figure(figsize=(8, 8))
plt.plot(fpr, tpr, label=f'ROC Curve (AUC = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], 'k--', label='Random Guessing')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend()
plt.show()

```



```

In [ ]: # Interpret results of the confusion matrix
if tp + fp == 0 or tp + fn == 0:
    print("Unable to determine if the model is overfitting or underfitting.")
else:
    precision_positive = tp / (tp + fp)
    recall_positive = tp / (tp + fn)
    specificity = tn / (tn + fp)
    sensitivity = recall_positive

```

```

balance = 1 - abs(1 - (precision_positive + sensitivity) / 2)

print("Precision (Positive):", precision_positive)
print("Recall (Positive):", recall_positive)
print("Specificity:", specificity)
print("Sensitivity:", sensitivity)
print("Balance:", balance)

if balance > 0.99:
    print("The model may be overfitting.")
elif balance < 0.05:
    print("The model may be underfitting.")
else:
    print("The model is reasonably balanced.")

```

```

Precision (Positive): 0.952191235059761
Recall (Positive): 0.9579158316633266
Specificity: 0.952191235059761
Sensitivity: 0.9579158316633266
Balance: 0.9550535333615437
The model is reasonably balanced.

```

```

In [ ]: # Initialize Decision Tree model with gini criterion (default)
dt_gini_model = DecisionTreeClassifier()

```

```

In [ ]: # Fit the model
dt_gini_model.fit(X_train, y_train)

```

```

Out[ ]: ▾ DecisionTreeClassifier
DecisionTreeClassifier()

```

```

In [ ]: # Predict on the test set
y_pred_dt_gini = dt_gini_model.predict(X_test)

```

```

In [ ]: # Evaluate the model
accuracy_dt_gini = accuracy_score(y_test, y_pred_dt_gini)
print("Decision Tree (Gini) Accuracy:", accuracy_dt_gini)

```

```

Decision Tree (Gini) Accuracy: 0.952047952047952

```

```

In [ ]: # Additional evaluation metrics
print("Classification Report:")
print(classification_report(y_test, y_pred_dt_gini))

```

```

Classification Report:

```

	precision	recall	f1-score	support
0	0.95	0.96	0.95	502
1	0.96	0.95	0.95	499
accuracy			0.95	1001
macro avg	0.95	0.95	0.95	1001
weighted avg	0.95	0.95	0.95	1001

```

In [ ]: classification_reports['Desicion Tree Normal'] = classification_report(y_test, y_pred_dt_gini)

```

```

In [ ]: # K-fold cross-validation and average accuracy
cv_accuracy = cross_val_score(dt_gini_model, X, y_bin, cv=10, scoring='accuracy')
avg_cv_accuracy = cv_accuracy.mean()
print("Average Cross-Validation Accuracy:", avg_cv_accuracy)

```

```

Average Cross-Validation Accuracy: 0.9644115768463072

```

```

In [ ]: # Confusion Matrix
conf_matrix = confusion_matrix(y_test, y_pred_dt_gini)
print("Confusion Matrix:")
print(conf_matrix)

```

```

Confusion Matrix:
[[480  22]
 [ 26 473]]

```

```

In [ ]: # Accuracy, Error rate, Precision, Recall, F-measure
tp, fp, fn, tn = conf_matrix.ravel()
accuracy = (tp + tn) / (tp + fp + fn + tn)
error_rate = 1 - accuracy
precision = tp / (tp + fp)
recall = tp / (tp + fn)
f_measure = 2 * (precision * recall) / (precision + recall)

```

```

In [ ]: print("Accuracy:", accuracy)

```

```

print("Error Rate:", error_rate)
print("Precision:", precision)
print("Recall:", recall)
print("F-measure:", f_measure)

```

```

Accuracy: 0.952047952047952
Error Rate: 0.047952047952047994
Precision: 0.9561752988047809
Recall: 0.9486166007905138
F-measure: 0.9523809523809524

```

```

In [ ]: # ROC Curve
y_scores = cross_val_predict(dt_gini_model, X, y_bin.ravel(), cv=10, method="predict_proba")[:, 1]
fpr, tpr, thresholds = roc_curve(y_bin.ravel(), y_scores)
roc_auc = auc(fpr, tpr)

print("ROC AUC:", roc_auc)


```

```
ROC AUC: 0.9642625349860056
```

```

In [ ]: # Visualize ROC Curve
plt.figure(figsize=(8, 8))
plt.plot(fpr, tpr, label=f'ROC Curve (AUC = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], 'k--', label='Random Guessing')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend()
plt.show()

```



```

In [ ]: # Interpret results of the confusion matrix
if tp + fp == 0 or tp + fn == 0:
    print("Unable to determine if the model is overfitting or underfitting.")
else:
    precision_positive = tp / (tp + fp)
    recall_positive = tp / (tp + fn)
    specificity = tn / (tn + fp)
    sensitivity = recall_positive
    balance = 1 - abs(1 - (precision_positive + sensitivity) / 2)

    print("Precision (Positive):", precision_positive)
    print("Recall (Positive):", recall_positive)

```

```

print("Specificity:", specificity)
print("Sensitivity:", sensitivity)
print("Balance:", balance)

if balance > 0.99:
    print("The model may be overfitting.")
elif balance < 0.05:
    print("The model may be underfitting.")
else:
    print("The model is reasonably balanced.")

```

Precision (Positive): 0.9561752988047809  
 Recall (Positive): 0.9486166007905138  
 Specificity: 0.9555555555555556  
 Sensitivity: 0.9486166007905138  
 Balance: 0.9523959497976473  
 The model is reasonably balanced.

#### Linear Discriminant Analysis (LDA):

```

In [ ]: from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
        from sklearn.model_selection import train_test_split
        from sklearn.metrics import accuracy_score, classification_report

```

```

In [ ]: X = df.drop('gender', axis=1)
        y = df['gender']

```

```

In [ ]: # Convert categorical labels to binary labels
        lb = LabelBinarizer()
        y_bin = lb.fit_transform(y)

```

```

In [ ]: X_train, X_test, y_train, y_test = train_test_split(X, y_bin, test_size=0.2, random_state=42)

```

```

In [ ]: # Initialize LDA model
        lda_model = LinearDiscriminantAnalysis()

```

```

In [ ]: # Fit the model
        lda_model.fit(X_train, y_train)

```

/opt/conda/lib/python3.10/site-packages/sklearn/utils/validation.py:1143: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples, ), for example using ravel().  
 y = column\_or\_1d(y, warn=True)

```

Out[ ]: ▾ LinearDiscriminantAnalysis
        LinearDiscriminantAnalysis()

```

```

In [ ]: # Predict on the test set
        y_pred_lda = lda_model.predict(X_test)

```

```

In [ ]: # Evaluate the model
        accuracy_lda = accuracy_score(y_test, y_pred_lda)
        print("LDA Accuracy:", accuracy_lda)

```

LDA Accuracy: 0.9590409590409591

```

In [ ]: # Additional evaluation metrics
        print("Classification Report:")
        print(classification_report(y_test, y_pred_lda))

```

Classification Report:

	precision	recall	f1-score	support
0	0.95	0.97	0.96	502
1	0.97	0.95	0.96	499
accuracy			0.96	1001
macro avg	0.96	0.96	0.96	1001
weighted avg	0.96	0.96	0.96	1001

```

In [ ]: classification_reports['LDA'] = classification_report(y_test, y_pred_lda)

```

```

In [ ]: # K-fold cross-validation and average accuracy
        cv_accuracy = cross_val_score(lda_model, X, y_bin, cv=10, scoring='accuracy')
        avg_cv_accuracy = cv_accuracy.mean()
        print("Average Cross-Validation Accuracy:", avg_cv_accuracy)

```

Average Cross-Validation Accuracy: 0.9694099800399201

```

/opt/conda/lib/python3.10/site-packages/sklearn/utils/validation.py:1143: DataConversionWarning: A column-vector
y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel()
l().
y = column_or_1d(y, warn=True)
/opt/conda/lib/python3.10/site-packages/sklearn/utils/validation.py:1143: DataConversionWarning: A column-vector
y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel()
l().
y = column_or_1d(y, warn=True)
/opt/conda/lib/python3.10/site-packages/sklearn/utils/validation.py:1143: DataConversionWarning: A column-vector
y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel()
l().
y = column_or_1d(y, warn=True)
/opt/conda/lib/python3.10/site-packages/sklearn/utils/validation.py:1143: DataConversionWarning: A column-vector
y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel()
l().
y = column_or_1d(y, warn=True)
/opt/conda/lib/python3.10/site-packages/sklearn/utils/validation.py:1143: DataConversionWarning: A column-vector
y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel()
l().
y = column_or_1d(y, warn=True)
/opt/conda/lib/python3.10/site-packages/sklearn/utils/validation.py:1143: DataConversionWarning: A column-vector
y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel()
l().
y = column_or_1d(y, warn=True)
/opt/conda/lib/python3.10/site-packages/sklearn/utils/validation.py:1143: DataConversionWarning: A column-vector
y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel()
l().
y = column_or_1d(y, warn=True)
/opt/conda/lib/python3.10/site-packages/sklearn/utils/validation.py:1143: DataConversionWarning: A column-vector
y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel()
l().
y = column_or_1d(y, warn=True)
/opt/conda/lib/python3.10/site-packages/sklearn/utils/validation.py:1143: DataConversionWarning: A column-vector
y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel()
l().
y = column_or_1d(y, warn=True)

```

```

In [ ]: # Confusion Matrix
conf_matrix = confusion_matrix(y_test, y_pred_lda)
print("Confusion Matrix:")
print(conf_matrix)

```

Confusion Matrix:

```

[[486  16]
 [ 25 474]]

```

```

In [ ]: # Accuracy, Error rate, Precision, Recall, F-measure
tp, fp, fn, tn = conf_matrix.ravel()
accuracy = (tp + tn) / (tp + fp + fn + tn)
error_rate = 1 - accuracy
precision = tp / (tp + fp)
recall = tp / (tp + fn)
f_measure = 2 * (precision * recall) / (precision + recall)

print("Accuracy:", accuracy)
print("Error Rate:", error_rate)
print("Precision:", precision)
print("Recall:", recall)
print("F-measure:", f_measure)

```

Accuracy: 0.9590409590409591  
Error Rate: 0.040959040959040904  
Precision: 0.9681274900398407  
Recall: 0.9510763209393346  
F-measure: 0.9595261599210265

```

In [ ]: # ROC Curve
y_scores = cross_val_predict(lda_model, X, y_bin.ravel(), cv=10, method="predict_proba")[:, 1]
fpr, tpr, thresholds = roc_curve(y_bin.ravel(), y_scores)
roc_auc = auc(fpr, tpr)

print("ROC AUC:", roc_auc)

```

ROC AUC: 0.9962665333866453

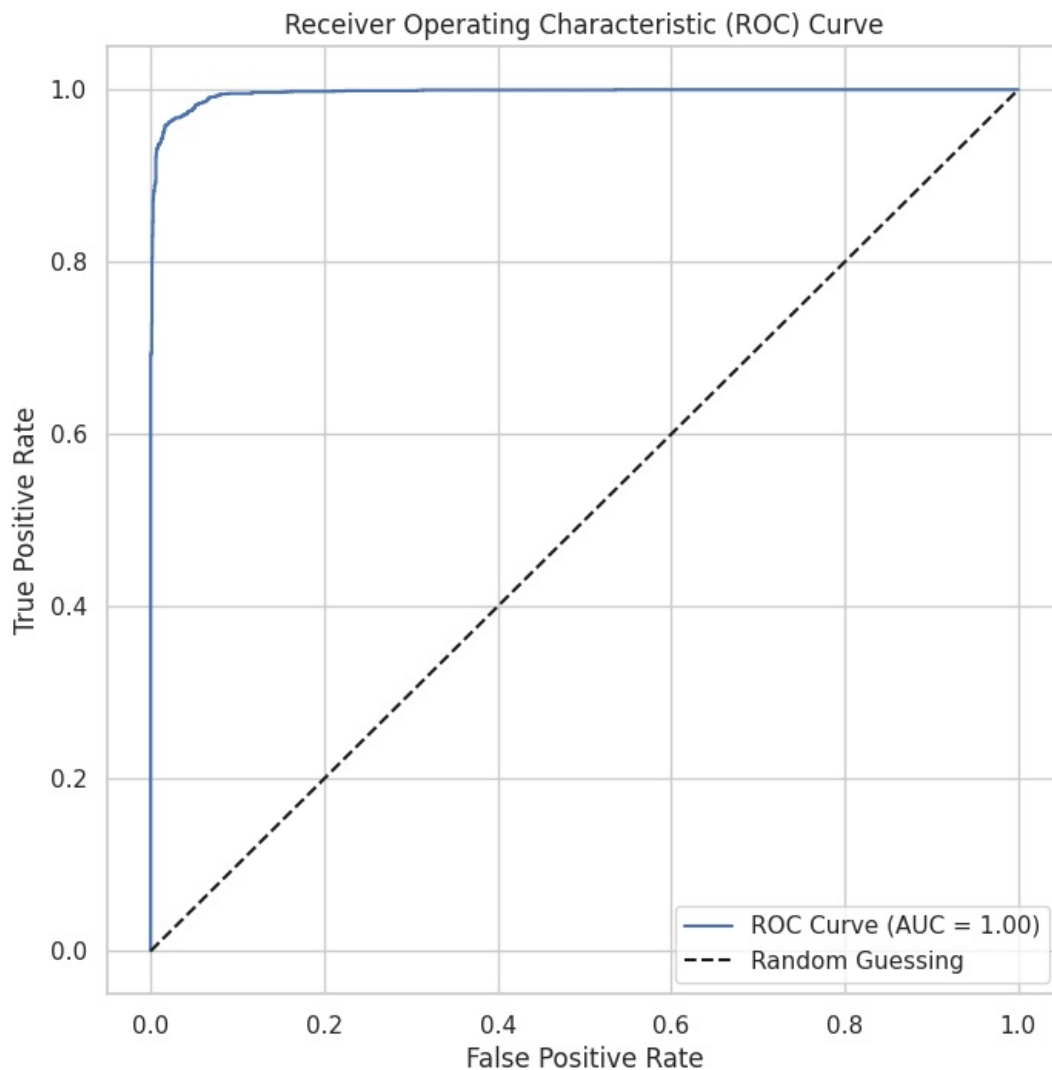
```

In [ ]: # Visualize ROC Curve
plt.figure(figsize=(8, 8))
plt.plot(fpr, tpr, label=f'ROC Curve (AUC = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], 'k--', label='Random Guessing')
plt.title('Receiver Operating Characteristic (ROC) Curve')

```



```
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend()
plt.show()
```



```
In [ ]: # Interpret results of the confusion matrix
if tp + fp == 0 or tp + fn == 0:
    print("Unable to determine if the model is overfitting or underfitting.")
else:
    precision_positive = tp / (tp + fp)
    recall_positive = tp / (tp + fn)
    specificity = tn / (tn + fp)
    sensitivity = recall_positive
    balance = 1 - abs(1 - (precision_positive + sensitivity) / 2)

    print("Precision (Positive):", precision_positive)
    print("Recall (Positive):", recall_positive)
    print("Specificity:", specificity)
    print("Sensitivity:", sensitivity)
    print("Balance:", balance)

    if balance > 0.99:
        print("The model may be overfitting.")
    elif balance < 0.05:
        print("The model may be underfitting.")
    else:
        print("The model is reasonably balanced.")
```

```
Precision (Positive): 0.9681274900398407
Recall (Positive): 0.9510763209393346
Specificity: 0.9673469387755103
Sensitivity: 0.9510763209393346
Balance: 0.9596019054895877
The model is reasonably balanced.
```

### Neural Network

```
In [ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import StratifiedKFold
```

```

from sklearn.metrics import roc_curve, auc, accuracy_score, confusion_matrix, precision_score, recall_score, f1
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from sklearn.model_selection import train_test_split

```

```

In [ ]: X = df.drop('gender', axis=1)
        y = df['gender']

```

```

In [ ]: y_bin = (y == 'Male').astype(int)

```

```

In [ ]: # Train-test split
        X_train, X_test, y_train, y_test = train_test_split(X, y_bin, test_size=0.2, random_state=42)

```

```

In [ ]: # Convert features and labels to numpy arrays
        X_train_np = X_train.values
        X_test_np = X_test.values
        y_train_np = y_train.values
        y_test_np = y_test.values

```

```

In [ ]: # Neural Network Model
        def create_model():
            model = Sequential()
            model.add(Dense(64, activation='relu'))
            model.add(Dense(1, activation='sigmoid'))
            model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
            return model

```

```

In [ ]: # K-fold Cross Validation for Neural Network
        k_fold = StratifiedKFold(n_splits=10, shuffle=True, random_state=42)
        cv_scores = []

```

```

In [ ]: # Lists to store evaluation metrics
        conf_matrices = []
        accuracies = []
        errors = []
        precisions = []
        recalls = []
        f1_scores = []

        for train_idx, test_idx in k_fold.split(X_train_np, y_train_np):
            # Initialize the model
            model = create_model()

            # Train the model on the current fold
            model.fit(X_train_np[train_idx], y_train_np[train_idx], epochs=10, batch_size=32, verbose=0)

            # Make predictions on the test set
            y_pred_nn = np.round(model.predict(X_train_np[test_idx])).astype(int)

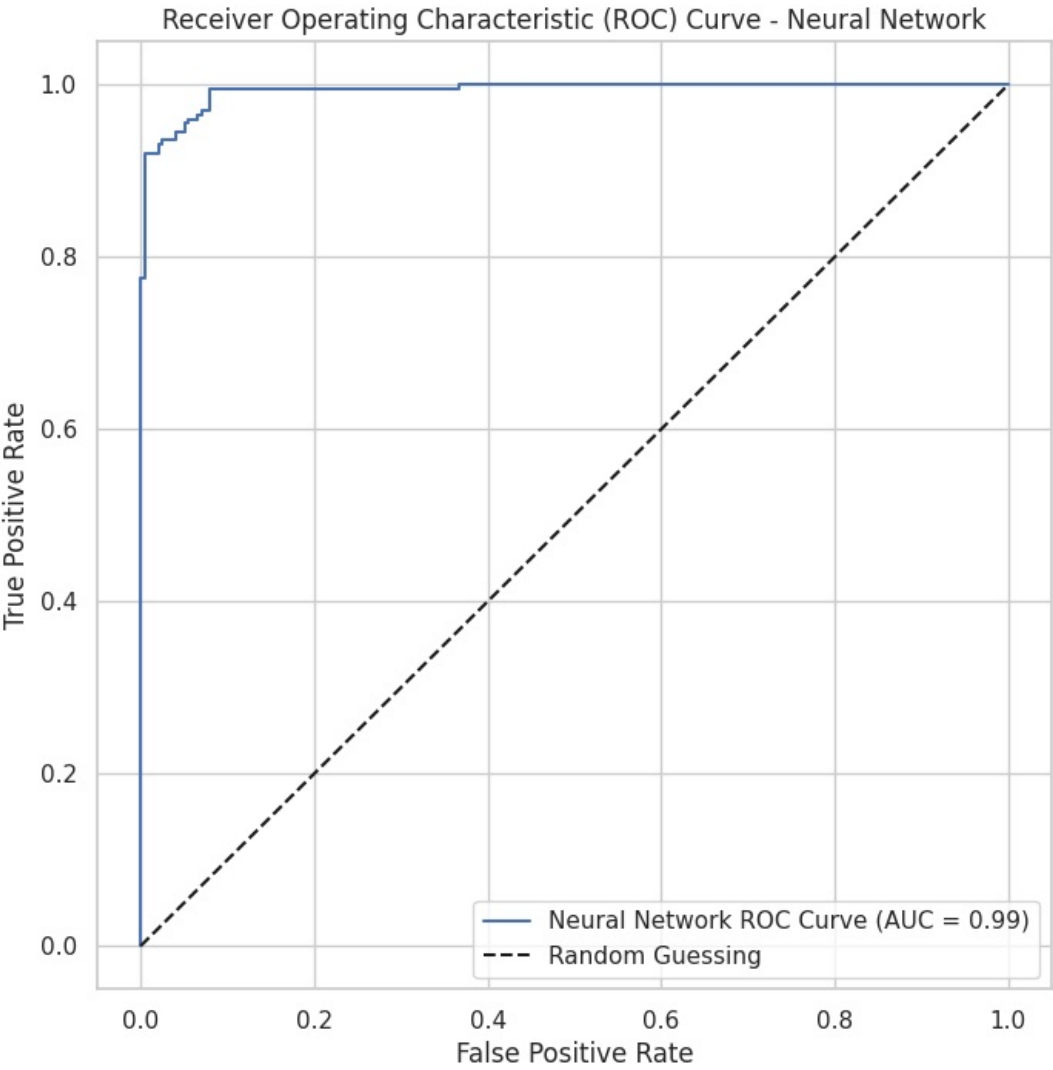
            # Evaluation metrics
            conf_matrix = confusion_matrix(y_train_np[test_idx], y_pred_nn)
            accuracy = accuracy_score(y_train_np[test_idx], y_pred_nn)
            error = 1 - accuracy
            precision = precision_score(y_train_np[test_idx], y_pred_nn)
            recall = recall_score(y_train_np[test_idx], y_pred_nn)
            f1 = f1_score(y_train_np[test_idx], y_pred_nn)


            # Append metrics to lists
            conf_matrices.append(conf_matrix)
            accuracies.append(accuracy)
            errors.append(error)
            precisions.append(precision)
            recalls.append(recall)
            f1_scores.append(f1)

            # Print cv_scores within the loop
            print("CV Scores:", cv_scores)


            # Visualize ROC Curve
            plt.figure(figsize=(8, 8))
            y_scores_nn = model.predict(X_train_np[test_idx])
            fpr_nn, tpr_nn, _ = roc_curve(y_train_np[test_idx], y_scores_nn)
            roc_auc_nn = auc(fpr_nn, tpr_nn)
            plt.plot(fpr_nn, tpr_nn, label=f'Neural Network ROC Curve (AUC = {roc_auc_nn:.2f})')
            plt.plot([0, 1], [0, 1], 'k--', label='Random Guessing')
            plt.title('Receiver Operating Characteristic (ROC) Curve - Neural Network')
            plt.xlabel('False Positive Rate')
            plt.ylabel('True Positive Rate')
            plt.legend()
            plt.show()

```






13/13 [=====] - 0s 1ms/step  
CV Scores: []  
13/13 [=====] - 0s 1ms/step

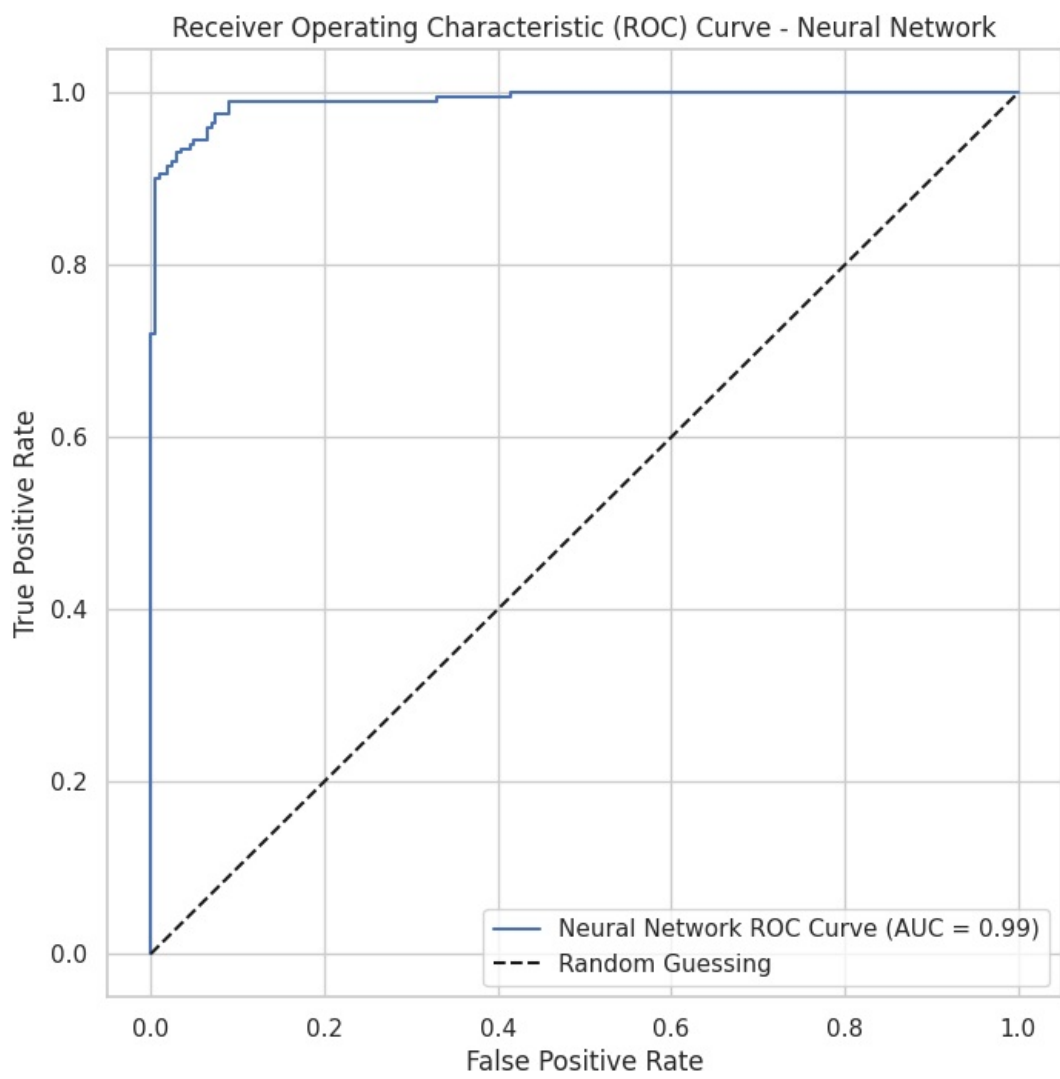


13/13 [=====] - 0s 1ms/step  
CV Scores: []  
13/13 [=====] - 0s 1ms/step




13/13 [=====] - 0s 1ms/step  
CV Scores: []  
13/13 [=====] - 0s 1ms/step







13/13 [=====] - 0s 1ms/step  
CV Scores: []  
13/13 [=====] - 0s 1ms/step




13/13 [=====] - 0s 1ms/step  
CV Scores: []  
13/13 [=====] - 0s 1ms/step




13/13 [=====] - 0s 1ms/step  
CV Scores: []  
13/13 [=====] - 0s 1ms/step



13/13 [=====] - 0s 1ms/step  
CV Scores: []  
13/13 [=====] - 0s 1ms/step



13/13 [=====] - 0s 2ms/step  
CV Scores: []  
13/13 [=====] - 0s 1ms/step



```
In [ ]: # Calculate average accuracy
if len(accuracies) > 0:
    average_accuracy = sum(accuracies) / len(accuracies)
    print("\nK-fold Cross Validation - Neural Network:")
    print("Average Accuracy:", average_accuracy)
else:
    print("\nNo accuracies to calculate average.")
```

K-fold Cross Validation - Neural Network:  
Average Accuracy: 0.95175

```
In [ ]: y_pred_nn = np.round(model.predict(X_test)).astype(int)
accuracy = accuracy_score(y_test, y_pred_nn)
print(f"Accuracy on the test set: {accuracy * 100:.2f}%")
```

32/32 [=====] - 0s 1ms/step  
Accuracy on the test set: 95.30%

```
In [ ]: print("Classification Report:")
print(classification_report(y_test, y_pred_nn))
```

Classification Report:

	precision	recall	f1-score	support
0	0.99	0.92	0.95	502
1	0.92	0.99	0.95	499
accuracy			0.95	1001
macro avg	0.95	0.95	0.95	1001
weighted avg	0.96	0.95	0.95	1001

```
In [ ]: classification_reports['NN'] = classification_report(y_test, y_pred_nn)
```

```
In [ ]: # Interpret results of the confusion matrix
if tp + fp == 0 or tp + fn == 0:
    print("Unable to determine if the model is overfitting or underfitting.")
else:
    precision_positive = tp / (tp + fp)
    recall_positive = tp / (tp + fn)
    specificity = tn / (tn + fp)
```



```

sensitivity = recall_positive
balance = 1 - abs(1 - (precision_positive + sensitivity) / 2)

print("Precision (Positive):", precision_positive)
print("Recall (Positive):", recall_positive)
print("Specificity:", specificity)
print("Sensitivity:", sensitivity)
print("Balance:", balance)

if balance > 0.99:
    print("The model may be overfitting.")
elif balance < 0.05:
    print("The model may be underfitting.")
else:
    print("The model is reasonably balanced.")

```

```

Precision (Positive): 0.9681274900398407
Recall (Positive): 0.9510763209393346
Specificity: 0.9673469387755103
Sensitivity: 0.9510763209393346
Balance: 0.9596019054895877
The model is reasonably balanced.

```

#### k-Nearest Neighbors (k-NN) with Different Distances:

```

In [ ]: from sklearn.neighbors import KNeighborsClassifier
        from sklearn.model_selection import train_test_split
        from sklearn.metrics import accuracy_score, classification_report

```

```

In [ ]: X = df.drop('gender', axis=1)
        y = df['gender']

```

```

In [ ]: y_bin = (y == 'Male').astype(int)

```

```

In [ ]: X_train, X_test, y_train, y_test = train_test_split(X, y_bin, test_size=0.2, random_state=42)

```

```

In [ ]: # Initialize k-NN model with Euclidean distance
        knn_euclidean_model = KNeighborsClassifier(n_neighbors=3, metric='euclidean')

```

```

In [ ]: # Fit the model
        knn_euclidean_model.fit(X_train, y_train)

```

```

Out[ ]: KNeighborsClassifier
        KNeighborsClassifier(metric='euclidean', n_neighbors=3)

```

```

In [ ]: # Predict on the test set
        y_pred_knn_euclidean = knn_euclidean_model.predict(X_test)

```

```

In [ ]: # Evaluate the model
        accuracy_knn_euclidean = accuracy_score(y_test, y_pred_knn_euclidean)
        print("k-NN (Euclidean Distance) Accuracy:", accuracy_knn_euclidean)

```

```
k-NN (Euclidean Distance) Accuracy: 0.964035964035964
```

```

In [ ]: # Additional evaluation metrics
        print("Classification Report:")
        print(classification_report(y_test, y_pred_knn_euclidean))

```

```

Classification Report:

```

	precision	recall	f1-score	support
0	0.95	0.98	0.96	502
1	0.98	0.95	0.96	499
accuracy			0.96	1001
macro avg	0.96	0.96	0.96	1001
weighted avg	0.96	0.96	0.96	1001

```

In [ ]: classification_reports['Knn Euc'] = classification_report(y_test, y_pred_knn_euclidean)

```

```

In [ ]: # K-fold cross-validation and average accuracy
        cv_accuracy = cross_val_score(knn_euclidean_model, X, y_bin, cv=10, scoring='accuracy')
        avg_cv_accuracy = cv_accuracy.mean()
        print("Average Cross-Validation Accuracy:", avg_cv_accuracy)

```

```
Average Cross-Validation Accuracy: 0.9674095808383232
```

```

In [ ]: # Confusion Matrix
        conf_matrix = confusion_matrix(y_test, y_pred_knn_euclidean)
        print("Confusion Matrix:")

```

```
print(conf_matrix)
```

Confusion Matrix:

```
[[490  12]
 [ 24 475]]
```

```
In [ ]: # Accuracy, Error rate, Precision, Recall, F-measure
tp, fp, fn, tn = conf_matrix.ravel()
accuracy = (tp + tn) / (tp + fp + fn + tn)
error_rate = 1 - accuracy
precision = tp / (tp + fp)
recall = tp / (tp + fn)
f_measure = 2 * (precision * recall) / (precision + recall)

print("Accuracy:", accuracy)
print("Error Rate:", error_rate)
print("Precision:", precision)
print("Recall:", recall)
print("F-measure:", f_measure)
```


```
Accuracy: 0.964035964035964
Error Rate: 0.03596403596403597
Precision: 0.9760956175298805
Recall: 0.953307392996109
F-measure: 0.9645669291338582
```

```
In [ ]: # ROC Curve
y_scores = cross_val_predict(knn_euclidean_model, X, y_bin.ravel(), cv=10, method="predict_proba")[:, 1]
fpr, tpr, thresholds = roc_curve(y_bin.ravel(), y_scores)
roc_auc = auc(fpr, tpr)

print("ROC AUC:", roc_auc)
```

```
ROC AUC: 0.9845687325069972
```

```
In [ ]: # Visualize ROC Curve
plt.figure(figsize=(8, 8))
plt.plot(fpr, tpr, label=f'ROC Curve (AUC = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], 'k--', label='Random Guessing')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend()
plt.show()
```



```
In [ ]: # Interpret results of the confusion matrix
if tp + fp == 0 or tp + fn == 0:
    print("Unable to determine if the model is overfitting or underfitting.")
else:
    precision_positive = tp / (tp + fp)
    recall_positive = tp / (tp + fn)
    specificity = tn / (tn + fp)
    sensitivity = recall_positive
    balance = 1 - abs(1 - (precision_positive + sensitivity) / 2)

    print("Precision (Positive):", precision_positive)
    print("Recall (Positive):", recall_positive)
    print("Specificity:", specificity)
    print("Sensitivity:", sensitivity)
    print("Balance:", balance)

    if balance > 0.99:
        print("The model may be overfitting.")
    elif balance < 0.05:
        print("The model may be underfitting.")
    else:
        print("The model is reasonably balanced.")
```

```
Precision (Positive): 0.9760956175298805
Recall (Positive): 0.953307392996109
Specificity: 0.9753593429158111
Sensitivity: 0.953307392996109
Balance: 0.9647015052629948
The model is reasonably balanced.
```

```
In [ ]: # Initialize k-NN model with Manhattan distance
knn_manhattan_model = KNeighborsClassifier(n_neighbors=3, metric='manhattan')
```

```
In [ ]: # Fit the model
knn_manhattan_model.fit(X_train, y_train)
```

```
Out[ ]: KNeighborsClassifier
KNeighborsClassifier(metric='manhattan', n_neighbors=3)
```

```
In [ ]: # Predict on the test set
y_pred_knn_manhattan = knn_manhattan_model.predict(X_test)
```

```
In [ ]: # Evaluate the model
accuracy_knn_manhattan = accuracy_score(y_test, y_pred_knn_manhattan)
print("k-NN (Manhattan Distance) Accuracy:", accuracy_knn_manhattan)
```

k-NN (Manhattan Distance) Accuracy: 0.9630369630369631

```
In [ ]: # Additional evaluation metrics
print("Classification Report:")
print(classification_report(y_test, y_pred_knn_manhattan))
```

Classification Report:

	precision	recall	f1-score	support
0	0.95	0.98	0.96	502
1	0.98	0.95	0.96	499
accuracy			0.96	1001
macro avg	0.96	0.96	0.96	1001
weighted avg	0.96	0.96	0.96	1001

```
In [ ]: classification_reports['Knn maha'] = classification_report(y_test, y_pred_knn_manhattan)
```

```
In [ ]: # K-fold cross-validation and average accuracy
cv_accuracy = cross_val_score(knn_manhattan_model, X, y_bin, cv=10, scoring='accuracy')
avg_cv_accuracy = cv_accuracy.mean()
print("Average Cross-Validation Accuracy:", avg_cv_accuracy)
```

Average Cross-Validation Accuracy: 0.9664095808383234

```
In [ ]: # Confusion Matrix
conf_matrix = confusion_matrix(y_test, y_pred_knn_manhattan)
print("Confusion Matrix:")
print(conf_matrix)
```

Confusion Matrix:  
[[490 12]  
 [ 25 474]]

```
In [ ]: # Accuracy, Error rate, Precision, Recall, F-measure
tp, fp, fn, tn = conf_matrix.ravel()
accuracy = (tp + tn) / (tp + fp + fn + tn)
error_rate = 1 - accuracy
precision = tp / (tp + fp)
recall = tp / (tp + fn)
f_measure = 2 * (precision * recall) / (precision + recall)

print("Accuracy:", accuracy)
print("Error Rate:", error_rate)
print("Precision:", precision)
print("Recall:", recall)
print("F-measure:", f_measure)
```


Accuracy: 0.9630369630369631  
Error Rate: 0.03696303696303693  
Precision: 0.9760956175298805  
Recall: 0.9514563106796117  
F-measure: 0.9636184857423796

```
In [ ]: # ROC Curve
y_scores = cross_val_predict(nb_model, X, y_bin.ravel(), cv=10, method="predict_proba")[:, 1]
fpr, tpr, thresholds = roc_curve(y_bin.ravel(), y_scores)
roc_auc = auc(fpr, tpr)

print("ROC AUC:", roc_auc)
```

ROC AUC: 0.9966251099560175

```
In [ ]: # Visualize ROC Curve
plt.figure(figsize=(8, 8))
plt.plot(fpr, tpr, label=f'ROC Curve (AUC = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], 'k--', label='Random Guessing')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend()
plt.show()
```



```
In [ ]: # Interpret results of the confusion matrix
if tp + fp == 0 or tp + fn == 0:
    print("Unable to determine if the model is overfitting or underfitting.")
else:
    precision_positive = tp / (tp + fp)
    recall_positive = tp / (tp + fn)
    specificity = tn / (tn + fp)
    sensitivity = recall_positive
    balance = 1 - abs(1 - (precision_positive + sensitivity) / 2)

    print("Precision (Positive):", precision_positive)
    print("Recall (Positive):", recall_positive)
    print("Specificity:", specificity)
    print("Sensitivity:", sensitivity)
    print("Balance:", balance)

    if balance > 0.99:
        print("The model may be overfitting.")
    elif balance < 0.05:
        print("The model may be underfitting.")
    else:
        print("The model is reasonably balanced.")
```

```
Precision (Positive): 0.9760956175298805
Recall (Positive): 0.9514563106796117
Specificity: 0.9753086419753086
Sensitivity: 0.9514563106796117
Balance: 0.9637759641047461
The model is reasonably balanced.
```

**Comparisons with Other Related Work on the Same Domain:**

```
In [ ]: # Print or access the reports later
for model, report in classification_reports.items():
    print(f"Classification Report for {model}:")
    print(report)
```

Classification Report for Naive Bayes:

	precision	recall	f1-score	support
0	0.96	0.97	0.96	502
1	0.97	0.96	0.96	499
accuracy			0.96	1001
macro avg	0.96	0.96	0.96	1001
weighted avg	0.96	0.96	0.96	1001

Classification Report for Desicion Tree Entropy:

	precision	recall	f1-score	support
0	0.96	0.95	0.96	502
1	0.95	0.96	0.96	499
accuracy			0.96	1001
macro avg	0.96	0.96	0.96	1001
weighted avg	0.96	0.96	0.96	1001

Classification Report for Desicion Tree Normal:

	precision	recall	f1-score	support
0	0.95	0.96	0.95	502
1	0.96	0.95	0.95	499
accuracy			0.95	1001
macro avg	0.95	0.95	0.95	1001
weighted avg	0.95	0.95	0.95	1001

Classification Report for LDA:

	precision	recall	f1-score	support
0	0.95	0.97	0.96	502
1	0.97	0.95	0.96	499
accuracy			0.96	1001
macro avg	0.96	0.96	0.96	1001
weighted avg	0.96	0.96	0.96	1001

Classification Report for NN:

	precision	recall	f1-score	support
0	0.99	0.92	0.95	502
1	0.92	0.99	0.95	499
accuracy			0.95	1001
macro avg	0.95	0.95	0.95	1001
weighted avg	0.96	0.95	0.95	1001

Classification Report for Knn Euc:

	precision	recall	f1-score	support
0	0.95	0.98	0.96	502
1	0.98	0.95	0.96	499
accuracy			0.96	1001
macro avg	0.96	0.96	0.96	1001
weighted avg	0.96	0.96	0.96	1001

Classification Report for Knn maha:

	precision	recall	f1-score	support
0	0.95	0.98	0.96	502
1	0.98	0.95	0.96	499
accuracy			0.96	1001
macro avg	0.96	0.96	0.96	1001
weighted avg	0.96	0.96	0.96	1001

## References (Papers Using the Same Data Sets):

[Link](#)